
COMP2017 / COMP9017

Week 8 Tutorial

Introduction to processes

Introduction to processes

- Helpful bash programs - `size`, `readelf` or `nm`, `objdump`, `fg`, `bg`

File descriptors

In week 6, we encountered file descriptors. Part of processes and soon, interprocess communication will involve using file descriptors. When opening or creating a file (or even opening a process) we will engage with file descriptors. Commonly referred to as `fd`.

Memory layout of a process

A process memory layout is typically made of (from lowest memory address to highest):

- Program text (Program instructions!) (`.text`)
- Initialised Global and Static Variables (Data Segment) (`.data`)
- Uninitialised global and static variables (`.bss`)
- Heap Memory (from low to high)
- Memory mapped `.data` and `.text` regions (Shared)
- Stack Memory (from high to low)
- Environment and args

And if you want to examine this you can use the `readelf` command that will allow you to examine the section headers of your executable.

```
readelf --section <your compiled program>
```

You can examine your executable by using the program `objdump` to grab the disassembly of `.text`

```
objdump -D -M intel <your compiled program>
```

Output from `hello_world`

```
...
40050d:      c7 45 fc 04 00 00 00 mov    DWORD PTR [rbp-0x4],0x4
400514:      c7 45 ec 00 00 00 00 mov    DWORD PTR [rbp-0x14],0x0
40051b:      eb 13                      jmp    400530 <main+0x43>
...
```

Using `fork()`

The `fork` function is used when creating a new child process. This function will duplicate the current calling process and its current execution context and allocate separate memory. Pointers and variables that are written to will eventually be written to a new memory space via CoW (Copy-on-Write). A new pid is generated for the sub process that can be isolated from the parent process.

Quote from man pages:

```
fork()  creates  a  new  process  by  duplicating  the  calling  process.  The
        new process is referred to as the child process.
        The calling process is referred to as the parent process.
```

Fork usage example:

```
#include<stdio.h>
#include<unistd.h>
#include<sys/wait.h>
#include<sys/types.h>

int main(void) {

    int n = 4;
    puts("about to fork");
    pid_t pid = fork();

    if (pid < 0) {
        perror("unable to fork");
        return 1;
    }

    if (pid == 0) {
        puts("child");
        n++;
        // sleep(1);
    } else {
        puts("parent");
    }
}
```

```
    n *= 2;
    // wait(NULL);
}

printf("fork returned %d, n is %d\n", pid, n);
return 0;
}
```

wait () and waitpid

The `wait` or `waitpid` system call will halt a process until another process has finished executing. The `wait` function will return the `pid` value of the child process.

```
pid_t pid = fork();
...

wait(NULL);
//alternatively waitpid(pid, NULL, 0);
```

The `wait` will implicitly wait for one of its children finishes execution while `waitpid` will explicitly wait for a process to finish with a specified `pid`.

Question 1: Signal Talk between Parent and Child

Within this task you will need to write an executable that will communicate between parent and a child. The parent will know the `pid` of the child and you will need to also ensure that the child knows its parents `pid`.

Once a `fork` has been executed, the two processes will communicate via signals.

The parent program should handle signals received:

- `SIGINT` -> The child wants you to focus on it (Send `SIGINT` back)

The child program should handle these signals received:

- `SIGINT` -> You will have the focus from the parent, yell: "I'm AMAZING!"
- `SIGHUP` -> Reply back to the parent that it will go outside and play (Shutdown the process)

Please look at the `man` pages for `waitpid`, `kill` and `pause`

exec () family calls

The `exec` family of functions allow a process to start an executable file. When a process calls `exec`, it is replaced in memory by the new process. If we want to start a new process and keep the current

one running, we need to call `fork` before calling `exec`. The following code uses `exec` to invoke the `sort` program located in `/usr/bin/sort` to sort the lines in the previous program, `forkdemo.c`. `exec` is given the path to a program to execute, then the contents of `argv[]` terminated by a null pointer. There are many variations of `exec` documented in the man pages.

```
#include <stdio.h>
#include <unistd.h>

int main(void) {
    puts("about to launch sort");

    if (execl("/usr/bin/sort", "sort", "forkdemo.c", NULL) == -1) {
        perror("exec failed");
        return 1;
    }

    puts("finished sort");
    return 0;
}
```

Question 2: Miniterminal

Observing how your current shell is able to execute a program by splitting it off into a separate process and executing. Try and recreate the same application using `fork` and `exec`.

Usage:

```
<command/program> <arguments>
```

Your miniterminal application will need to pass in command line arguments when initialised. Your program terminal program should parse the command line and split them by the spaces as defined and be fault tolerant.

Example:

```
cat file.txt
```

Your terminal should await for user input like so:

```
>
```

For this task you will need to use `fork` and `exec`. Once a command has been inputted by the user, your terminal must fork the current process and then proceed to call `exec` with the correct command. It must maintain focus on the process it just executed.

For the sake of simplicity, you can assume a few number of arguments.

Question 3: Controlling Processes

After you have written a simple miniterminal that can execute a processes, we will extend this to maintain a list or (some data structure of your choosing) of running processes and using the `waitpid` function, after a command is executed you can check to see if all processes that have been executed by miniterminal are still running.

If a process is no longer running, it is the job of your terminal to clean remove the process from your table.

This kind of functionality is aimed at maintaining background processes when using `&`.

For example: Example:

```
cat file.txt &
```

If the input provided by the user contains an `&`, your program should not provide any focus for this process.

You can observe this behaviour in your own terminal window by executing a process with `&`.

It is advised you read `waitpid` man page for information on how to retrieve the status without blocking.

popen and pclose

In a similar manner to `open` and `close` we are able to run a command and in turn we are provided a `FILE*` object that we can read from or write to (not both however). We can retrieve output from `stdout` from a process and pipe it into our own program.

```
FILE* f = popen("<program>", )
...
fread(...)
...

//Close it in similar fashion
pclose(f);
```

Question 4: Log output from a program

Using `popen` and `pclose` create a process logger. A process logger will launch a program and extract the output from that program to a file specified.

```
./proc_log --cmd="python3 say_hello_100.py" --out=test.out
```

With the provided command and out flags, your program should execute the program with the specified command and pipe its output to a file.

Something fun: Use the socat command and create a channel between from your process logger and use the cat command to read the other end.

Example use for socat

```
socat pty,raw,echo=0,link=./program_input pty,raw,echo=0,link=./program_out
```

Question 5: Reading a linux executable (Extension)

You are tasked with reading the data of program executable, this program can be the program you are currently writing to read a program executable.

Specification of a linux executable

```
typedef struct elf64_hdr {
    unsigned char    e_ident[EI_NIDENT];
    Elf64_Half e_type;
    Elf64_Half e_machine;
    Elf64_Word e_version;
    Elf64_Addr e_entry;
    Elf64_Off e_phoff;
    Elf64_Off e_shoff;
    Elf64_Word e_flags;
    Elf64_Half e_ehsize;
    Elf64_Half e_phentsize;
    Elf64_Half e_phnum;
    Elf64_Half e_shentsize;
    Elf64_Half e_shnum;
    Elf64_Half e_shstrndx;
} Elf64_Ehdr;
```

Question 6: Break your compiled program (Extension)

Using nm or readelf and objdump, you will change the regularly printed out value of your code to be something entirely different.

Constructing a hello_number application you can change the register or static variable's data so it will print out the number 7.

References

Chapter 9, Page 839-843