

---

# COMP2017 / COMP9017      Week 10 Tutorial

---

## Parallelism with POSIX threads and optimisation in C

### Pthreads basics

Pthreads is a library specified by the POSIX standard, defining an API for creating and manipulating threads.

To use the Pthreads library:

1. Include the `pthread.h` header in your source code.
2. Specify `-pthread` in your compiler flags to tell the compiler to link and configure the Pthreads library.

```
> clang -g -Wall -Werror -std=gnu11 program.c -o program -pthread
```

Linker flags should be specified at the end of your usual compiler options.

### Creating threads

To create a thread using the Pthreads library, we use the `pthread_create` function.

```
int pthread_create(  
    pthread_t* thread,  
    const pthread_attr_t* attr,  
    void* (*start_routine) (void *),  
    void* arg  
);
```

1. The first argument to this function takes a `pthread_t`, it will use it to store the thread ID.
2. The second argument is a pointer to a `pthread_attr_t` structure, you can use this to specify additional options in the creation of the thread, it's fine to leave this as `NULL`.
3. The third argument is a pointer to a function that the thread will execute once it spawns.
4. The last argument is the parameter passed to the thread function.
5. The function returns a non zero value if an error occurred.

## Waiting on threads

The `pthread_join` function forces the calling thread to wait for a particular thread ID to finish.

```
int pthread_join(pthread_t thread, void** retval);
```

1. The first argument is the thread to wait for.
2. The second argument allows you to get the return value from the `start_routine`. If you don't need the return value you can just pass in `NULL`.
3. The function returns a non-zero value if an error occurred.

## Question 1: Hello from Pthreads

The following code creates 4 threads, giving each of them an argument and waits for them to finish.

```
#include <stdio.h>
#include <pthread.h>
#define NTHREADS 4
void* worker(void* arg) {
    const int argument = *((int*) arg);
    printf("Hello from thread %d\n", argument);
    return NULL;
}

int main(void) {
    int args[NTHREADS] = { 1, 2, 3, 4 };
    pthread_t thread_ids[NTHREADS];
    for (size_t i = 0; i < NTHREADS; i++) { // Create threads with worker fn
        if (pthread_create(thread_ids + i, NULL, worker, args + i) != 0) {
            perror("unable to create thread");
            return 1;
        }
    }
    for (size_t i = 0; i < NTHREADS; i++) { // Wait for all threads to finish
        if (pthread_join(thread_ids[i], NULL) != 0) {
            perror("unable to join thread");
            return 1;
        }
    }
}
```

Compile and run the code above. Make sure you understand exactly what each line is doing. Reminder that `printf` is thread safe by the POSIX standard, thus calling it from multiple threads will yield no side effects.

1. Why does the program output Hello from thread ... in different orderings?
2. What is the difference between a thread (`pthread_create`) and a process (`fork`)?

## Question 2: Parallel Sum

In the following code, each thread is given a section of an array of numbers to sum.

1. How long does it take to run when you set `THREADS` to 1, 2, 3 ... etc?  
You can use `time ./sum` to see how long the program takes to run.
2. Why does using more threads make the program slower?
3. How can this be fixed? Hint: [False sharing](#)

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

#define LENGTH 100000000
#define NTHREADS 4
#define NREPEATS 10
#define CHUNK (LENGTH / NTHREADS)

typedef struct {
    size_t id;
    long* array;
    long result;
} worker_args;

void* worker(void* args) {

    worker_args* wargs = (worker_args*) args;

    const size_t start = wargs->id * CHUNK;
    const size_t end = wargs->id == NTHREADS - 1 ? LENGTH : (wargs->id + 1) * C

    // Sum values from start to end
    for (size_t i = start; i < end; i++) {
        wargs->result += wargs->array[i];
    }

    return NULL;
}
```

```
int main(void) {

    long* numbers = malloc(sizeof(long) * LENGTH);
    for (size_t i = 0; i < LENGTH; i++) {
        numbers[i] = i + 1;
    }

    worker_args* args = malloc(sizeof(worker_args) * NTHREADS);
    for (size_t n = 1; n <= NREPEATS; n++) {
        for (size_t i = 0; i < NTHREADS; i++) {
            args[i] = (worker_args) {
                .id      = i,
                .array   = numbers,
                .result   = 0,
            };
        }

        pthread_t thread_ids[NTHREADS];

        // Launch threads
        for (size_t i = 0; i < NTHREADS; i++) {
            pthread_create(thread_ids + i, NULL, worker, args + i);
        }

        // Wait for threads to finish
        for (size_t i = 0; i < NTHREADS; i++) {
            pthread_join(thread_ids[i], NULL);
        }

        long sum = 0;

        // Calculate total sum
        for (size_t i = 0; i < NTHREADS; i++) {
            sum += args[i].result;
        }

        printf("Run %2zu: total sum is %ld\n", n, sum);
    }

    free(args);
    free(numbers);
}
```

## Question 3: Performance and benchmarking

The execution time of a program can be measured by using the `time` command. Taking the example from the next exercise, we can measure how long it takes to run by running the `time` command like so:

```
time ./mutex
4000000

real    0m0.420s
user    0m0.999s
sys     0m1.234s
```

The *real* time is the wall clock time taken to run the program. The *user* time is the total CPU time used, and the *system* time is the time spent in system calls or the kernel. Here you can see that the CPU time spent is larger than the wall clock time spent, so this was executing on multiple cores simultaneously.

Sometimes it is better to have finer-grained measures of time, for example if your program has a large setup phase, and you only want to measure some operation which occurs after that. For this you can use the `clock()` function. This provides you with microsecond accuracy. You can achieve nanosecond accuracy using by using operating system specific [functions](#).

```
#include <time.h>
#include <stdio.h>

int main(void) {

    const clock_t tick = clock();

    int ops = 0;
    for (int i = 0; i < 10000000; i++) {
        ops += i;
    }

    const clock_t tock = clock();
    printf("Time elapsed: %fs\n", (double) (tock - tick) / CLOCKS_PER_SEC);

    return 0;
}
```

## Question 4: Matrix multiplication

The code below performs matrix multiplication.

1. Analyse the program's performance:

To use `gprof`, you need to compile with profiling using `gcc` and the `-pg` compiler flag.

```
$ gcc -g -pg -Wall -Werror -std=gnull matrix.c -o matrix -pthread
```

- (a) Ensure that the `-pg` flag has been used to compile the program.
  - (b) Execute the program as usual.
  - (c) After the program has finished, execute the profiler:  

```
$ gprof ./matrix > mm.stats
```
  - (d) View the saved information using `less`, which function is the bottleneck in the code and why?
  - (e) You can use `gprof -l` to show statistics for each line of the source.
2. Swap the order of the loops so that it is `y, k, x` instead of `y, x, k`.  
Why does this have any effect on the performance of the code?
  3. Improve the performance by using `pthread`s for parallelism.
  4. Measure how long it takes to run with and without parallelism using `time`. Recommend that you use the lab machines to benchmark as they are similar to the machines used for the upcoming assignments. When benchmarking, it's important that no other programs are running since they may affect the results.
  5. Compare your performance with the cache optimised version in [What every programmer should know about memory by Ulrich Drepper](#).

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <assert.h>
#include <pthread.h>

#define WIDTH 512
#define IDX(x, y) ((y) * WIDTH + (x))

/**
 * Returns the matrix multiplication of a and b.
 */
float* multiply(const float* a, const float* b) {

    float* result = calloc(WIDTH * WIDTH, sizeof(float));

    for (size_t y = 0; y < WIDTH; y++) {
        for (size_t x = 0; x < WIDTH; x++) {
            for (size_t k = 0; k < WIDTH; k++) {
                result[IDX(x, y)] += a[IDX(k, y)] * b[IDX(x, k)];
            }
        }
    }

    return result;
}

/**
 * Returns a Hadamard matrix, if H is Hadamard matrix, then
 *  $HH^T = nI$ , where  $I$  is the identity matrix and  $n$  is the width.
 * Easy to verify that the matrix multiplication was done correctly.
 *
 * Sylvester's construction implemented here only works
 * for matrices that have width that is a power of 2.
 *
 * Note that this construction produces matrices that are symmetric.
 */
float* hadamard(void) {

    // Ensure the width is a power of 2
    assert(((WIDTH - 1) & WIDTH) == 0);

    size_t w = WIDTH;
    size_t quad_size = 1;

    float* result = malloc(WIDTH * WIDTH * sizeof(float));

    result[0] = 1;
    while ((w >= 1) != 0) {
        // Duplicate the upper left quadrant into the other three quadrants
    }
}
```

```
    for (size_t y = 0; y < quad_size; ++y) {
        for (size_t x = 0; x < quad_size; ++x) {
            const float v = result[IDX(x, y)];
            result[IDX(x + quad_size, y)] = v;
            result[IDX(x, y + quad_size)] = v;
            result[IDX(x + quad_size, y + quad_size)] = -v;
        }
    }

    quad_size *= 2;
}

return result;
}

// Displays a matrix.
void display(const float* matrix) {

    for (size_t y = 0; y < WIDTH; y++) {
        for (size_t x = 0; x < WIDTH; x++) {
            printf("%6.2f ", matrix[IDX(x, y)]);
        }
        printf("\n");
    }
}

int main(void) {

    // Construct the matrices
    float* a = hadamard();
    float* b = hadamard();

    // Compute the result
    float* r = multiply(a, b);

    // Verify the result
    for (size_t y = 0; y < WIDTH; y++) {
        for (size_t x = 0; x < WIDTH; x++) {
            assert(x == y ? r[IDX(x, y)] == WIDTH : r[IDX(x, y)] == 0);
        }
    }

    puts("done");

    free(a);
    free(b);
    free(r);

    return 0;
}
```



## Question 5: Critical sections and mutual exclusion

A critical section is a piece of code that accesses a shared resource that must not be concurrently accessed by more than one thread. These are usually enforced using locks or other synchronisation measures.

Mutexes are a way of enforcing a critical section. In our threads, we can call `pthread_mutex_lock` in order to obtain ownership of the mutex before entering the critical section. If the thread cannot obtain ownership of the mutex object, it will wait until the mutex has been unlocked. The thread that owns the mutex must call `pthread_mutex_unlock` after it has exited the critical section.

The example below shows a typical use of mutexes in pthreads. This mutex is statically initialised using `PTHREAD_MUTEX_INITIALIZER`. You can also create them dynamically using `pthread_mutex_init`.

```
#include <stdio.h>
#include <pthread.h>
#define THREADS 4
#define LOOPS 1000000

static unsigned counter = 0;
static pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;

static void* worker(void *arg) {

    for (unsigned i = 0; i < LOOPS; i++) {
        // attempt to lock the mutex...
        // thread will wait when mutex is already locked
        pthread_mutex_lock(&mutex);
        // only one thread will be able execute this code
        counter += 1;
        // unlock the mutex after the critical section
        pthread_mutex_unlock(&mutex);
    }

    return NULL;
}

int main(void) {

    pthread_t thread_ids[THREADS];
    for (size_t i = 0; i < THREADS; i++) {
        pthread_create(thread_ids + i, NULL, worker, NULL);
    }

    for (size_t i = 0; i < THREADS; i++) {
        pthread_join(thread_ids[i], NULL);
    }
    printf("%d\n", counter);
}
```

1. Run the code and verify that it outputs the result correctly.
2. Comment out the `pthread_mutex_unlock` line, does the program behave in a way you expect?
3. Uncomment out both `pthread_mutex_lock` and `pthread_mutex_unlock` lines, repeat the program until it outputs an incorrect result, why does this happen?